


Image Filtering

EECS 442 – David Fouhey

Winter 2023, University of Michigan

https://web.eecs.umich.edu/~fouhey/teaching/EECS442_W23/

Let's Take An Image



Let's Fix Things

- We have noise in our image
- Let's replace each pixel with a *weighted* average of its neighborhood
- Weights are *filter kernel*

	Out	


1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

1D Case

Signal




Filter



What's the average
of 9, 10, 12?

Output



(a) 9

(b) 11.5

(c) 10.33


(d) 11.66

1D Case


Signal



Filter



Output



Done! Next?


1D Case

(1) 10.66 (2) 9.33
(3) 14.2 (4) 11.33


Signal



Filter



Output




1D Case

- (1) 10.33
- (2) 11.33
- (3) 10
- (4) 9.1


Signal



Filter




Output



1D Case


Signal



Filter




Output



1D Case


Signal



Filter



Output



1D Case

10	12	9	11	10	11	12
----	----	---	----	----	----	----

*

$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
---------------	---------------	---------------

=

10.33	10.66	10	10.66	11
-------	-------	----	-------	----

You lose pixels (maybe...)
Filter “sees” only a few pixels at a time

Applying a Linear Filter

Input

I11	I12	I13	I14	I15	I16
I21	I22	I23	I24	I25	I26
I31	I32	I33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

O11	O12	O13	O14
O21	O22	O23	O24
O31	O32	O33	O34

Applying a Linear Filter

Input & Filter

F11	F12	F13	I14	I15	I16
F21	F22	F23	I24	I25	I26
F31	F32	F33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Output

O11

$$O11 = I11 * F11 + I12 * F12 + \dots + I33 * F33$$

Applying a Linear Filter

Input & Filter

I11	F11	F12	F13	I15	I16
I21	F21	F22	F23	I25	I26
I31	F31	F32	F33	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Output

O11	O12
-----	-----

$$O12 = I12 * F11 + I13 * F12 + \dots + I34 * F33$$

Applying a Linear Filter

Input

I11	I12	I13	I14	I15	I16
I21	I22	I23	I24	I25	I26
I31	I32	I33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

**How many times can we apply a
3x3 filter to a 5x6 image?**

Applying a Linear Filter

Input

I11	I12	I13	I14	I15	I16
I21	I22	I23	I24	I25	I26
I31	I32	I33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

O11	O12	O13	O14
O21	O22	O23	O24
O31	O32	O33	O34

$$O_{ij} = I_{ij} * F_{11} + I_{i(j+1)} * F_{12} + \dots + I_{i(j+2)} * F_{33}$$


Painful Details – Edge Cases

Filtering doesn't keep the whole image.

Suppose **f** is the image and **g** the filter.


Full – any part of **g** touches **f**. **Same** – same size as **f**;

Valid – only when filter doesn't fall off edge.



Painful Details – Edge Cases

What to about the “?” region?



Symm: fold sides over



Circular/Wrap: wrap around



pad/fill: add value, often 0




Painful Details – Does it Matter?

(I've applied the filter per-color channel)

Which padding did I use and why?

Input
Image



Filtered
???



Filtered
???




Note – this is a zoom of the filtered, not a filter of the zoomed

Painful Details – Does it Matter?

(I've applied the filter per-color channel)

Input
Image



Filtered
Symm Pad



Filtered
Zero Pad



Note – this is a zoom of the filtered, not a filter of the zoomed

Practice with Linear Filters



Original

0	0	0
0	1	0
0	0	0

?

Practice with Linear Filters



Original

0	0	0
0	1	0
0	0	0



The Same!

Practice with Linear Filters



Original

0	0	0
0	0	1
0	0	0

?

Practice with Linear Filters



Original

0	0	0
0	0	1
0	0	0



Shifted
LEFT
1 pixel

Practice with Linear Filters



Original

0	1	0
0	0	0
0	0	0

?

Practice with Linear Filters



Original

0	1	0
0	0	0
0	0	0



Shifted
DOWN
1 pixel

Practice with Linear Filters



Original

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$


?

Practice with Linear Filters



Original

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$



Blur
(Box Filter)

Practice with Linear Filters



Original

0	0	0
0	2	0
0	0	0

-

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

?

Practice with Linear Filters



Original

0	0	0
0	2	0
0	0	0


-

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$




Sharpened
(Accentuates
difference from
local average)

Sharpening



before



after

Properties – Linear

Assume: I image f1, f2 filters

Linear: $\text{apply}(I, f1+f2) = \text{apply}(I, f1) + \text{apply}(I, f2)$

I is a white box on black, and f1, f2 are rectangles

$$A\left(\begin{array}{|c|} \hline \blacksquare \\ \hline \end{array}, \begin{array}{|c|} \hline \blacksquare \\ \hline \end{array} + \begin{array}{|c|} \hline \square \\ \hline \end{array}\right) = A\left(\begin{array}{|c|} \hline \blacksquare \\ \hline \end{array}, \begin{array}{|c|} \hline \square \\ \hline \end{array}\right) = \begin{array}{|c|} \hline \text{blurred} \\ \hline \end{array}$$

$$A\left(\begin{array}{|c|} \hline \blacksquare \\ \hline \end{array}, \begin{array}{|c|} \hline \blacksquare \\ \hline \end{array}\right) + A\left(\begin{array}{|c|} \hline \blacksquare \\ \hline \end{array}, \begin{array}{|c|} \hline \square \\ \hline \end{array}\right) = \begin{array}{|c|} \hline \text{blurred} \\ \hline \end{array} + \begin{array}{|c|} \hline \text{blurred} \\ \hline \end{array} = \begin{array}{|c|} \hline \text{blurred} \\ \hline \end{array}$$



Note: I am showing filters un-normalized and blown up. They're a smaller box filter (i.e., each entry is $1/(\text{size}^2)$)

Properties – Shift-Invariant

Assume: I image, f filter

Shift-invariant: $\text{shift}(\text{apply}(I, f)) = \text{apply}(\text{shift}(I, f))$


Intuitively: only depends on filter neighborhood




Painful Details – Signal Processing

What I just showed is often called “convolution”.
Actually cross-correlation. Source of **terrible** confusion.

Cross-Correlation
(Original Orientation)



Convolution
(Flip filter in x,y first)




Convolution

To be more clear:

```
def imageFilter(image, filter):  
    for y in range(...): for x in range(...)  
        # you'll fill this in  
    return filtered  
  
def imageConvolve2D(image, filter):  
    # flip the filter left/right and up/down  
    filterUse = np.fliplr(np.flipud(filter))  
    return imageFilter2D(image, filterUse)
```

Properties of Convolution

- Any shift-invariant, linear operation is a convolution ($*$)
- Commutative: $f * g = g * f$
- Associative: $(f * g) * h = f * (g * h)$
- Distributes over $+$: $f * (g + h) = f * g + f * h$
- Scalars factor out: $kf * g = f * kg = k(f * g)$
- Identity (a single one with all zeros):





Questions?

- Nearly everything onwards is a convolution.
- This is important to get right.

Smoothing With A Box

Intuition: if filter touches it, it gets a contribution.



Input



Filter

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

Output



Solution – Weighted Combination



Intuition: weight according to closeness to center. Define 0,0 to be center pixel, then:

$$Filter_{ij} \propto 1$$

What's this?



$$Filter_{ij} \propto \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$




Recognize the Filter?

It's a Gaussian!

$$Filter_{ij} \propto \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$



0.003	0.013	0.022	0.013	0.003
0.013	0.060	0.098	0.060	0.013
0.022	0.098	0.162	0.098	0.022
0.013	0.060	0.098	0.060	0.013
0.003	0.013	0.022	0.013	0.003





Smoothing With A Box & Gauss

Still have some speckles, but it's not a big box


Input



Box Filter




Gauss. Filter



Gaussian Filters


$$\sigma = 1$$

filter = 21x21




$$\sigma = 2$$

filter = 21x21




$$\sigma = 4$$

filter = 21x21




$$\sigma = 8$$

filter = 21x21




Note: filter visualizations are independently normalized throughout the slides so you can see them better

Applying Gaussian Filters





Applying Gaussian Filters

Input Image
(no filter)





Applying Gaussian Filters

$$\sigma = 1$$




Applying Gaussian Filters

$$\sigma = 2$$





Applying Gaussian Filters

$\sigma = 4$



Applying Gaussian Filters

$\sigma = 8$




Picking a Filter Size

Too small filter \rightarrow bad approximation


Want size $\approx 6\sigma$ (99.7% of energy)

Left far too small; right slightly too small!

$\sigma = 8$, size = 21



$\sigma = 8$, size = 43



Runtime Complexity

Image size = $N \times N = 6 \times 6$

Filter size = $M \times M = 3 \times 3$

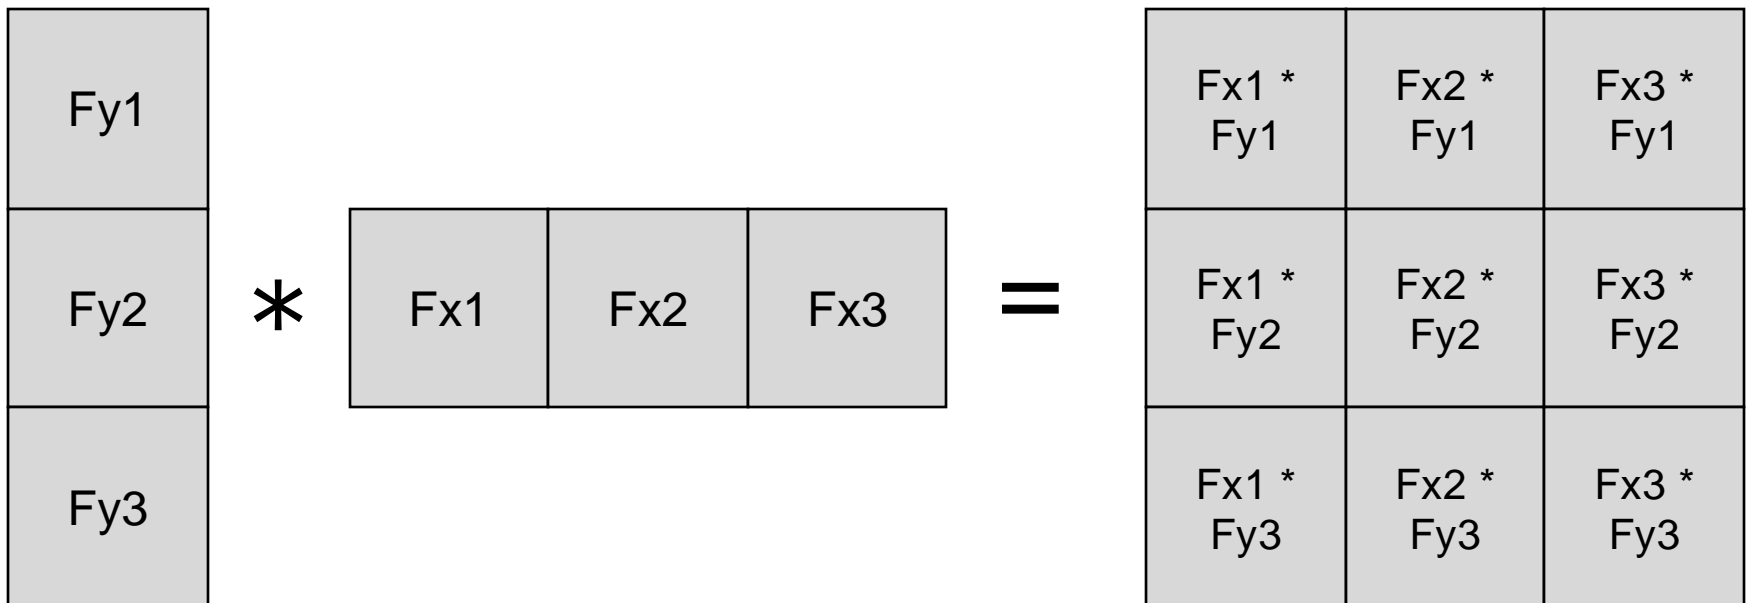
I11	I12	I13	I14	I15	I16
I21	F11	F12	F13	I25	I26
I31	F21	F22	F23	I35	I36
I41	F31	F32	F33	I45	I46
I51	I52	I53	I54	I55	I56
I61	I62	I63	I64	I65	I66

```
for ImageY in range(N):  
    for ImageX in range(N):  
        for FilterY in  
range(M):  
            for FilterX in  
range(M):  
                ...
```

Time: $O(N^2M^2)$

Separability

Conv(vector, transposed vector) \rightarrow outer product



Separability

$$Filter_{ij} \propto \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$


→

$$Filter_{ij} \propto \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right)$$

Separability

1D Gaussian * 1D Gaussian = 2D Gaussian

Image * 2D Gauss = Image * (1D Gauss * 1D Gauss)
= (Image * 1D Gauss) * 1D Gauss



Runtime Complexity

Image size = $N \times N = 6 \times 6$

Filter size = $M \times 1 = 3 \times 1$

I11	I12	I13	I14	I15	I16
I21	F1	I23	I24	I25	I26
I31	F2	I33	I34	I35	I36
I41	F3	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56
I61	I62	I63	I64	I65	I66

**What are my compute
savings for a 13x13 filter?**


```
for ImageY in range(N):  
    for ImageX in range(N):  
        for FilterY in  
range(M):  
        ...  
for ImageY in range(N):  
    for ImageX in range(N):  
        for FilterX in  
range(M):
```

Why Gaussian?

Gaussian filtering removes parts of the signal above a certain frequency. Often noise is high frequency and signal is low frequency.





Where Gaussian Fails



Applying Gaussian Filters

$$\sigma = 1$$



Why Does This Fail?

Means can be arbitrarily distorted by outliers

Signal

10	12	9	8	1000	11	10	12
----	----	---	---	------	----	----	----

Filter

0.1	0.8	0.1
-----	-----	-----

Output

11.5	9.2	107.3	801.9	109.8	10.3
------	-----	-------	-------	-------	------

What else is an “average” other than a mean?

Non-linear Filters (2D)

40	81	13	22
125	830	76	80
144	92	108	95
132	102	106	87

[040, 081, 013, 125, 830, 076, 144, 092, 108]

↓ Sort ↓

[013, 040, 076, 081, 092, 108, 125, 144, 830]

↓
92

[830, 076, 080, 092, 108, 095, 102, 106, 087]

↓ Sort ↓

[076, 080, 087, 092, 095, 102, 106, 108, 830]

↓
95



Applying Median Filter

Median
Filter
(size=3)




Applying Median Filter

Median
Filter
(size = 7)



Is Median Filtering Linear?

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 2 \\ 2 & 2 & 2 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 2 \\ 2 & 2 & 2 \end{bmatrix}$$



Filtering – Sharpening

Image




Smoothed



Details

=



Filtering – Sharpening

Image



$+\alpha$

Details



“Sharpened” $\alpha=1$

=



Filtering – Sharpening

Image



$+\alpha$

Details



“Sharpened” $\alpha=0$

$=$



Filtering – Sharpening

Image



$+\alpha$

Details



“Sharpened” $\alpha=2$

=



Filtering – Sharpening

Image



$+\alpha$

Details



“Sharpened” $\alpha=0$

=



Filtering – Extreme Sharpening


$$+\alpha$$


“Sharpened” $\alpha=10$



Filtering

What's this Filter?



-1	0	1
----	---	---

Dx

-1	0	1
----	---	---

^T

Dy



Filtering – Derivatives

$$(Dx^2 + Dy^2)^{1/2}$$



Filtering – Bonus

- If you're curious, you can use filters to accomplish a surprisingly large number of things.


Filtering – Counting

How many “on” pixels have
10+ neighbors within 10 pixels?

Pixels



Disk



*

=

???



Filtering – Counting


How many “on” pixels have
10+ neighbors within 10 pixels?

Pixels



X

Density



=

Answer

